NeurIPS 2024 Conference on Neural Information Processing Systems

Unified Feature Attribution and Counterfactual Explanations in Deep Learning

Abstract:
Explainable AI (XAI) is critical for trustworthy machine learning. In this work, we investigate interpretability mecha